

# Sentiment Analysis for Product Reviews

## 1. Main Objective

The primary goal of this analysis is to develop and compare deep learning models that classify customer product reviews into three sentiment classes: positive, neutral, or negative. By accurately detecting customer sentiment, businesses can rapidly address issues, enhance overall customer satisfaction, and make data-driven decisions on product improvements, marketing strategies, and support services.

## 2. Dataset Summary


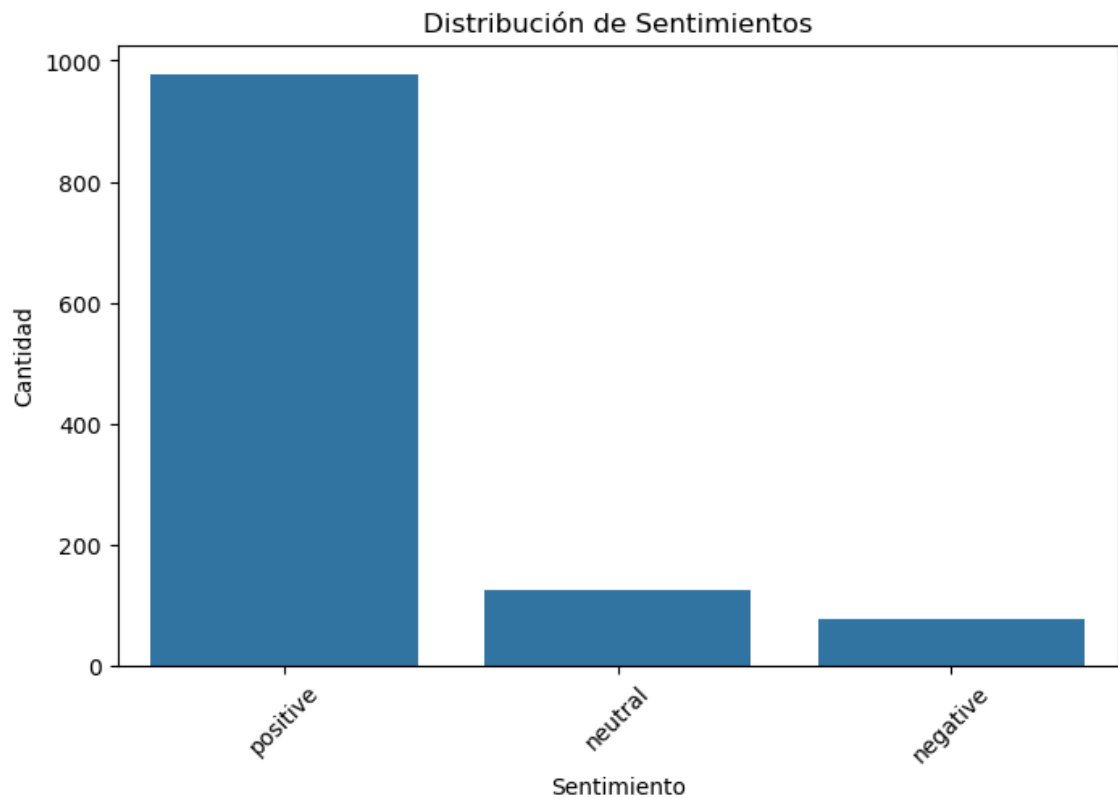
- **Data Source:** Amazon Product Reviews dataset, originally comprising 50,000 reviews.
- **Features:**
  - **Review text:** Main textual feedback provided by customers.
  - **Rating:** Numerical score provided by the customer.
  - **Sentiment label:** Derived from the rating (converted to 0, 1, or 2 for negative, neutral, or positive).
  - **Other metadata (optional):** Product category, review timestamp, and user information.

### Data Utilization:

- The **review text** is the central feature for building the sentiment classification models.
- **Ratings** were used to generate labels (0, 1, or 2) to supervise training.
- **Additional metadata** (if leveraged in further analysis) could reveal insights on specific product categories or seasonal trends in sentiment.

Before modeling, significant preprocessing steps were performed:

1. Text normalization (e.g., lowercasing, removing punctuation).
2. Removal of stop words and irrelevant characters.
3. Tokenization to split reviews into meaningful units.
4. Vectorization and embedding for deep learning input.



### 3. Deep Learning Models and Their Performance

Three deep learning approaches were tested to determine the most effective model for sentiment classification:

#### 1. LSTM (Long Short-Term Memory) Model

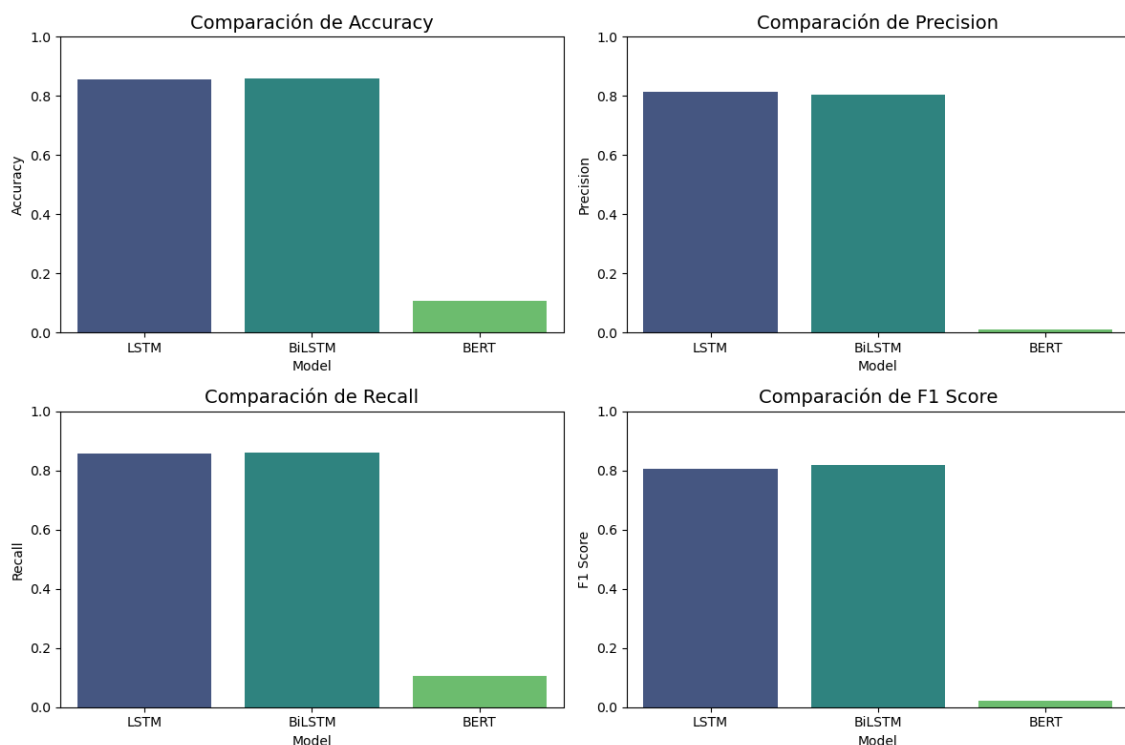
- **Summary:** Uses an LSTM layer to capture sequential dependencies in text.
- **Performance** (on test set):
  - Accuracy: ~85.59%
  - Weighted F1 Score: ~0.8052
  - Challenges handling very small classes (e.g., class 0 with only 15 samples).

## 2. Bidirectional LSTM Model

- **Summary:** Processes text in both forward and backward directions, improving context capture.
- **Performance:**
  - Accuracy: ~86.02%
  - Weighted F1 Score: ~0.8184
  - Shows better recall for the largest class (2), but still struggles with minority classes (0 and 1).

## 3. BERT (Transformer-Based) Model

- **Summary:** Fine-tuned from the pre-trained `bert-base-uncased` model, expected to capture more nuanced language.
- **Performance:**
  - Accuracy: ~10.59%
  - Weighted F1 Score: ~0.0203
  - The poor performance suggests a mismatch or potential issue in data encoding, hyperparameter configuration, or training steps. It may also be influenced by the extreme class imbalance.




## Model Best Suited to the Objective

- Despite BERT's typical superiority in many NLP tasks, the **Bidirectional LSTM** currently outperforms the other models given the reported metrics. The

**BiLSTM** shows the highest F1 score (~0.8184) and the best balance across classes under the current training setup.

- The BERT model's unexpectedly low performance indicates that additional debugging or adjustments are necessary (e.g., verifying tokenization, class weights, or re-checking the training pipeline).



#### 4. Key Findings and Observations

- **Class Imbalance:** The dataset appears highly skewed toward class 2 (positive reviews). Classes 0 (negative) and 1 (neutral) have significantly fewer samples, which negatively affects precision and recall metrics.
- **Minority Class Performance:** LSTM and BiLSTM models exhibit near-zero f1-scores for class 0, indicating difficulty in properly learning from scarce negative samples.
- **Model Behavior:**
  - Bidirectional LSTM handles context slightly better than the basic LSTM, improving performance in the largest class while moderately increasing recall for the neutral class.
  - BERT, which typically excels in complex language understanding, underperformed here, suggesting a possible flaw in how the data was encoded, how the model was trained, or how labels are distributed.

#### 5. Potential Flaws and Plan of Action

1. **Extreme Class Imbalance**
  - **Flaw:** Classes 0 and 1 have too few samples relative to class 2.
  - **Action:**

- Collect more data or apply oversampling/undersampling/augmentation techniques to balance classes.
  - Assign class weights in the loss function to penalize misclassification of minority classes.
- 2. **Data Encoding and Preprocessing for BERT**
  - **Flaw:** Very low accuracy for BERT suggests issues in the preprocessing or tokenization steps.
  - **Action:**
    - Revisit tokenization and ensure input shapes are correct.
    - Verify if the labels (0, 1, 2) align correctly with the model's output layer.
    - Explore learning rate and hyperparameter tuning specifically for BERT.
- 3. **Reviewing Hyperparameters**
  - **Flaw:** Default hyperparameters may not be optimal for a three-class problem with imbalanced data.
  - **Action:**
    - Conduct a systematic hyperparameter search (e.g., learning rate, batch size, epochs) for each model.
    - Investigate advanced architectures (e.g., ensembles) to improve minority class detection.
- 4. **Domain-Specific Adjustments**
  - **Flaw:** Generic preprocessing may overlook domain-specific terms or synonyms.
  - **Action:**
    - Incorporate a domain-specific vocabulary or additional textual features (e.g., product categories).
    - Fine-tune embeddings or BERT on a more domain-relevant corpus if available.

## 6. Next Steps

- **Data Augmentation:** Increase representation of minority sentiment classes (negative and neutral) to improve training.
- **Hyperparameter Tuning:** Use grid search or Bayesian optimization to find optimal configurations for learning rate, batch size, and sequence lengths.
- **Refine the BERT Pipeline:**
  - Confirm the correctness of input IDs, attention masks, and label mapping.
  - Verify the custom loss function (`categorical_crossentropy`) aligns with the label encoding.
  - Experiment with training BERT for more epochs, using a smaller batch size or a lower learning rate.
- **Ensemble Approaches:** Combine predictions from BiLSTM and a corrected BERT model to further improve robustness.
- **Periodic Retraining:** Integrate newly acquired reviews to keep the model updated, especially if product lines or customer behaviors change over time.

## 7. Questions to Refine the Analysis

To further enhance the report and your model performance, it would be helpful to clarify:

1. **Class Distribution in the Full Dataset:** How many samples belong to each of the three classes (negative, neutral, positive) before the train-test split?
2. **Preprocessing for BERT:** Could there be any mismatch between label encoding (0, 1, 2) and the BERT model's expectation of [0, 1, 2]?
3. **Loss Function:** Are you certain `categorical_crossentropy` is appropriate, or should we use `sparse_categorical_crossentropy` depending on how labels are encoded?
4. **Validation Data:** Are you using enough and balanced validation samples for monitoring the training process of BERT?
5. **Hyperparameter Settings:** Which specific optimizer parameters (learning rate, epsilon) and number of epochs did you use for BERT, and have you tried different settings?